

CS317

Information Retrieval

Week 07

Muhammad Rafi

February 26, 2024

Scoring, Term weighting &
Vector Space Model (VSM)

Chapter No. 6

Agenda

- Limitation of Boolean Retrieval Model
 - Ranked Retrieval /Scoring
 - Scoring between Query (Dq) and Documents (Di)
 - Zone Indexing & Learning Zone Weights
 - Vector space model
 - Conclusion
-

Boolean Model

- Boolean queries often result in either too few (=0) or too many (1000s) results.
 - Query 1: “*standard user dlink 650*” → 200,000 hits
 - Query 2: “*standard user dlink 650 no card found*”: 0 hits
 - It takes a lot of skill to come up with a query that produces a manageable number of hits.
 - AND gives too few; OR gives too many
-

Ranked Retrieval

- Rather than a set of documents satisfying a query expression, in ranked retrieval, the system returns an ordering over the (top) documents in the collection for a query
- Free text queries: Rather than a query language of operators and expressions, the user's query is just one or more words in a human language
- In principle, there are two separate choices here, but in practice, ranked retrieval has normally been associated with free text queries and vice versa

Score –ranking

- We wish to return in order the documents most likely to be useful to the searcher
- How can we rank-order the documents in the collection with respect to a query?
- Assign a score – say in $[0, 1]$ – to each document
- This score measures how well document and query “match”.

Query – Matching Documents

- We need a way of assigning a score to a query/document pair
- Let's start with a one-term query
- If the query term does not occur in the document: score should be 0
- The more frequent the query term in the document, the higher the score (should be)
- We will look at a number of alternatives for this.

First Attempt- Jaccard Coefficient

- Recall from Lecture 3: A commonly used measure of overlap of two sets A and B
- $\text{jaccard}(A,B) = |A \cap B| / |A \cup B|$
- $\text{jaccard}(A,A) = 1$
- $\text{jaccard}(A,B) = 0$ if $A \cap B = 0$
- A and B don't have to be the same size.
- Always assigns a number between 0 and 1.

Parametric Search

- Documents
 - Data
 - Metadata
- A parametric search provide IR based on Parameters.
- Example:
 - Find documents which contains “computational linguistics” in title, “Manning” in author and “Parameters” in the body of the text.

Example

Bibliographic Search

Search category	Value
Author	<i>Example: Widom, J or Garcia-Molina</i> <input type="text"/>
Title	<i>Also a part of the title possible</i> <input type="text"/>
Date of publication	<i>Example: 1997 or <1997 or >1997 limits the search to the documents appeared in, before and after 1997 respectively</i> <input type="text"/>
Language	Language the document was written in English ▾
Project	ANY ▾
Type	ANY ▾
Subject group	ANY ▾
Sorted by	Date of publication ▾

Weighted Zone Scoring

- Different fields/zones may have different importance in evaluating how a document matches a query
- For a query q and a document d , weighted zone scoring assigns to pair (q, d) a score in range $[0, 1]$ by computing a linear combination of zone scores
- Suppose each document has l zones, let $g_1, \dots, g_l \in [0, 1]$ such that $\sum_{i=1}^l g_i = 1$
 - Each field/zone of the document contributes a Boolean value – let s_i be the Boolean score denoting a match or absence between q and the i -th zone
 - The weighted zone score is $\sum_{i=1}^l g_i \times s_i$

Weighted Zone Indexing



Example 6.1: Consider the query *shakespeare* in a collection in which each document has three zones: *author*, *title*, and *body*. The Boolean score function for a zone takes on the value 1 if the query term *shakespeare* is present in the zone, and 0 otherwise. Weighted zone scoring in such a collection requires three weights g_1, g_2 , and g_3 , respectively corresponding to the *author*, *title*, and *body* zones. Suppose we set $g_1 = 0.2$, $g_2 = 0.3$, and $g_3 = 0.5$ (so that the three weights add up to 1); this corresponds to an application in which a match in the *author* zone is least important to the overall score, the *title* zone somewhat more, and the *body* contributes even more.

Thus, if the term *shakespeare* were to appear in the *title* and *body* zones but not the *author* zone of a document, the score of this document would be 0.8.

Learning Weights for Zone Scoring

- Using training examples that have been judged editorially
- Each training example is a tuple consisting of a query q and a document d , and a relevance judgment for d on q
 - The judgment can be binary – relevant or not
 - A judgment score can also be used
- Compute the weights such that the learned scores approximate the relevance judgments as much as possible
 - An optimization problem

Example – Learning Weights

Example	DocID	Query	s_T	s_B	Judgment
Φ_1	37	linux	1	1	Relevant
Φ_2	37	penguin	0	1	Non-relevant
Φ_3	238	system	0	1	Relevant
Φ_4	238	penguin	0	0	Non-relevant
Φ_5	1741	kernel	1	1	Relevant
Φ_6	2094	driver	0	1	Relevant
Φ_7	3191	driver	1	0	Non-relevant

Figure 6.5 An illustration of training examples.

Learning Weights for Zone Scoring

We now consider a simple case of weighted zone scoring, where each document has a *title* zone and a *body* zone. Given a query q and a document d , we use the given Boolean match function to compute Boolean variables $s_T(d, q)$ and $s_B(d, q)$, depending on whether the title (respectively, body) zone of d matches query q . For instance, the algorithm in Figure 6.4 uses an AND of the query terms for this Boolean function. We will compute a score between 0 and 1 for each (document, query) pair using $s_T(d, q)$ and $s_B(d, q)$ by using a constant $g \in [0, 1]$, as follows:

$$\text{score}(d, q) = g \cdot s_T(d, q) + (1 - g)s_B(d, q).$$

We now describe how to determine the constant g from a set of *training examples*, each of which is a triple of the form $\Phi_j = (d_j, q_j, r(d_j, q_j))$. In each training example, a given training document d_j and a given training query q_j are assessed by a human editor who delivers a relevance judgment $r(d_j, q_j)$ that is either *relevant* or *nonrelevant*. This is illustrated in Figure 6.5, where seven training examples are shown.

Learning Weights for Zone Scoring

For each training example Φ_j we have Boolean values $s_T(d_j, q_j)$ and $s_B(d_j, q_j)$ that we use to compute a score from (6.2)

$$\text{score}(d_j, q_j) = g \cdot s_T(d_j, q_j) + (1 - g)s_B(d_j, q_j).$$

We now compare this computed score with the human relevance judgment for the same document–query pair (d_j, q_j) ; to this end, we quantize each *relevant* judgment as a 1 and each *nonrelevant* judgment as a 0. Suppose that we define the error of the scoring function with weight g as

$$\varepsilon(g, \Phi_j) = (r(d_j, q_j) - \text{score}(d_j, q_j))^2,$$

where we have quantized the editorial relevance judgment $r(d_j, q_j)$ to 0 or 1. Then, the total error of a set of training examples is given by

$$\sum_j \varepsilon(g, \Phi_j).$$

Learning Weights for Zone Scoring

We begin by noting that for any training example Φ_j for which $s_T(d_j, q_j) = 0$ and $s_B(d_j, q_j) = 1$, the score computed by Equation (6.2) is $1 - g$. In similar fashion, we may write down the score computed by Equation (6.2) for the three other possible combinations of $s_T(d_j, q_j)$ and $s_B(d_j, q_j)$; this is summarized in Figure 6.6.

Let n_{01r} (respectively, n_{01n}) denote the number of training examples for which $s_T(d_j, q_j) = 0$ and $s_B(d_j, q_j) = 1$ and the editorial judgment is *relevant* (respectively, *nonrelevant*). Then the contribution to the total error in Equation (6.4) from training examples for which $s_T(d_j, q_j) = 0$ and $s_B(d_j, q_j) = 1$ is

$$[1 - (1 - g)]^2 n_{01r} + [0 - (1 - g)]^2 n_{01n}.$$

By writing in similar fashion the error contributions from training examples of the other three combinations of values for $s_T(d_j, q_j)$ and $s_B(d_j, q_j)$ (and extending the notation in the obvious manner), the total error corresponding to Equation (6.4) is

Learning Weights for Zone Scoring

$$(n_{01r} + n_{10n})g^2 + (n_{10r} + n_{01n})(1 - g)^2 + n_{00r} + n_{11n}.$$

By differentiating Equation (6.5) with respect to g and setting the result to 0, it follows that the optimal value of g is

$$\frac{n_{10r} + n_{01n}}{n_{10r} + n_{10n} + n_{01r} + n_{01n}}.$$

Bag of Words

- In this model, a text (such as a sentence or a document) is represented as the bag (multiset) of its words, disregarding grammar and even word order but keeping multiplicity.
 - A document is simply the collection of terms/words/lexemes with frequencies.
 - Similarly a query is also transformed as BoW representation.
-

Bag of Words

■ Example

D1: {John likes to watch movies. Mary likes movies too. }

D2: {John also likes to watch football games.}

D3: { Mary does not like football games }

■ Bow Representation

D1: {John, like, watch, movie, Mary, like, movie}

D2: {John, like, watch, football, game}

D3: { Mary, like, football, game }

■ Query Term

Q: {like football game}

Bag of Words

■ Scores

Common Terms in document and query over total terms

$$(D1, Q) = (D1 \cap Q) / (D1 \cup Q) = 1/8$$

$$(D2, Q) = 3/5$$

$$(D3, Q) = 3/4$$

■ Issues

Document size affects the overall similarity.

Frequency of common terms are ignored.

Similarity is only based on common terms, which are treated as independent of each other.

Vector Space Model is the answer to these problems.

Vector Space Model

- Vector Space Model is an algebraic model for representing text documents (and any objects, in general) as vectors of identifiers, such as, for example, terms of a document.
- It is used in information filtering, information retrieval, indexing and relevancy rankings.
- Its first use was in the SMART Information Retrieval System.

Vector Representation of Documents

- From offline collection build a dictionary.
- Let the size of vocabulary be “V”
- So a V-dimensional vector is used for representing documents in limited vocabulary space.
- This dictionary is kept in alphabetical order of vocabulary

$$d = \langle 0, 1, 0, 1, \dots, 0 \rangle$$

Vector Representation of Query

- The query is also represented similar to document.
 - A V-dimensional vector is also used for query, only the term present in query are non-zero vector for the query.
-

Cosine Similarity

- One of the limitation of Boolean Model is the similarity measure it offers.
- In vector space model, similarity is calculated by cosine measure between query and document.

$$\text{Cos } \Theta = (d \cdot q) / (|d| \cdot |q|)$$

Term Frequency

- In initial vector space model, we only use a term present for the v-dimensional calculation.
- What if a term appears more frequent in a document?
- Can we use term-frequency
 - Term Frequency – how many times a term appear in a document.
- Problem with frequent term?

Sub linear “ tf ” Scaling

- Term Frequency can be scaled in different ways (tf)
 - Log Linear $1 + \log(\text{tf}_{t,d})$
 - Augmented $0.5 + \frac{0.5 \times \text{tf}_{t,d}}{\max_t(\text{tf}_{t,d})}$
 - Maximum Frequency normalization
-

Document Frequency

- Document Frequency – How many documents contains a given term? (df)?
 - Scaling the df $\text{idf}_t = \log \frac{N}{\text{df}_t}$
 - Augmented $\max\{0, \log \frac{N - \text{df}_t}{\text{df}_t}\}$
-

Weighting Schemes

- One of the popular and effective weighting scheme is
 - term frequency * inverse document frequency

$$\text{tf-idf}_{t,d} = \text{tf}_{t,d} \times \text{idf}_t.$$

- In other words, tf-idf, assigns to term t a weight in document d that is
 - Highest when t occurs many times within a small number of documents (thus lending high discriminating power to those documents);
 - Lower when the term occurs fewer times in a document, or occurs in many documents (thus offering a less pronounced relevance signal);
 - Lowest when the term occurs in virtually all documents.

Different Variants

Term frequency	Document frequency	Normalization
n (natural) $\text{tf}_{t,d}$	n (no) 1	n (none) 1
l (logarithm) $1 + \log(\text{tf}_{t,d})$	t (idf) $\log \frac{N}{\text{df}_t}$	c (cosine) $\frac{1}{\sqrt{w_1^2 + w_2^2 + \dots + w_M^2}}$
a (augmented) $0.5 + \frac{0.5 \times \text{tf}_{t,d}}{\max_t(\text{tf}_{t,d})}$	p (prob idf) $\max\{0, \log \frac{N - \text{df}_t}{\text{df}_t}\}$	u (pivoted unique) $1/u$ (Section 6.4.4)
b (boolean) $\begin{cases} 1 & \text{if } \text{tf}_{t,d} > 0 \\ 0 & \text{otherwise} \end{cases}$		b (byte size) $1/\text{CharLength}^\alpha, \alpha < 1$
L (log ave) $\frac{1 + \log(\text{tf}_{t,d})}{1 + \log(\text{ave}_{t \in d}(\text{tf}_{t,d}))}$		

► **Figure 6.15** SMART notation for tf-idf variants. Here *CharLength* is the number of characters in the document.

Advantages of VSM

- Simple model based on linear algebra
 - Term weights not binary
 - Allows computing a continuous degree of similarity between queries and documents
 - Allows ranking documents according to their possible relevance
 - Allows partial matching
 - Last 20 years of working knowledge
-

Disadvantages of VSM

- The order in which the terms appear in the document is lost in the vector space representation.
 - Theoretically assumes terms are statistically independent.
 - Search keywords must precisely match document terms; word substrings might result in a "false positive match"
 - Semantic sensitivity; documents with similar context but different term vocabulary won't be associated, resulting in a "false negative match".
-

SMART System

- The SMART system did a rigor research on different schemes, for document and query they have used mnemonic ddd.qqq
- The first letter in each triplet specifies the term frequency component of the weighting, the second the document frequency component, and the third the form of normalization used.
- For example, a very standard weighting scheme is Inc.ltc, where the document vector has log-weighted term frequency, no idf (for both effectiveness and efficiency reasons), and cosine normalization, while the query vector uses log-weighted term frequency, idf weighting, and cosine normalization.